

OH, HEY, YOU ORGANIZED
OUR PHOTO ARCHIVE!

YEAH, I TRAINED A NEURAL
NET TO SORT THE UNLABELED
PHOTOS INTO CATEGORIES.

WHOA! NICE WORK!



ENGINEERING TIP:

WHEN YOU DO A TASK BY HAND,
YOU CAN TECHNICALLY SAY YOU
TRAINED A NEURAL NET TO DO IT.

INFO 251: Applied Machine Learning

Neural Networks, part 1

Course Outline

- Causal Inference and Research Design
 - Experimental methods
 - Non-experiment methods
- Machine Learning
 - Design of Machine Learning Experiments
 - Linear Models and Gradient Descent
 - Non-linear models
 - **Neural models**
 - Unsupervised Learning
 - Practicalities, Fairness, Bias
- Special topics

Key Concepts (last two lectures)

- Decision trees and regression trees
- Recursive tree algorithm
- Choosing splits
- Information gain
- Overfitting and pruning
- Regression trees
- Random forests
- AdaBoost
- Gradient boosting
- Feature Importance

Outline

- **Neural Networks: Motivation and Biology**
- The Perceptron
- Learning weights
- Multilayer networks
- Backpropagation
- Summary

Neural Networks

- Computational models inspired by the brain
 - Mimic how the brain processes information
 - In the hopes that computers can reason as well as human brains
 - And perhaps even better
 - ➔ Build machine learning algorithms based on the most sophisticated learner out there!

Engineering Brains

- What's a Brain?
 - Composed of 100 B neurons – we'll come back to this
 - Switching time: 0.001 seconds
 - *10,000 – 100,000 connections per neuron*
- Scene recognition
 - 0.1 seconds => Parallel computation
- Compare to transistor:
 - 100,000,000,000 transistors in modern chip (human $\times 10^3$)
 - *Switching time: 0.000000001 seconds* (human $\times 10^7$)
 - 10-100 connections per transistor

What's a Neuron?



What's a Neuron?





Outline

- Neural Networks: Motivation and Biology
- **The Perceptron**
- Learning weights
- Multilayer networks
- Backpropagation
- Summary

Creating Artificial Neurons


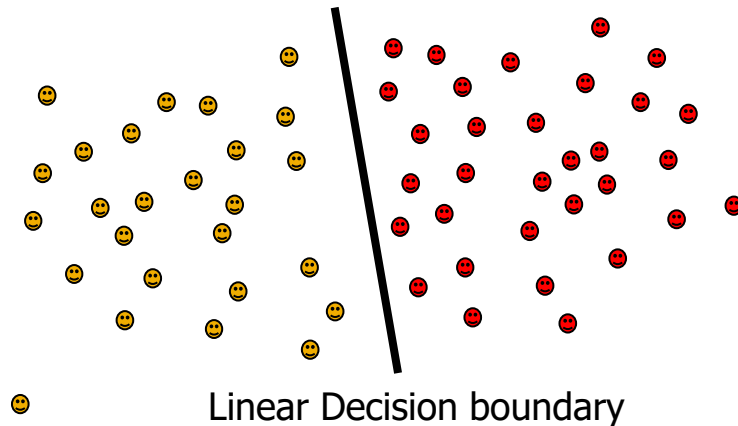
- How to model a neuron?
 - Neuron fires when membrane potential exceeds a threshold

- The "Perceptron"
 - Perceptron "fires" if sum of inputs exceeds threshold
 - $h(x) = \text{Sign}(b + \sum_{d=1}^k w_d x_d)$
 - k weights indexed by w_d
 - Bias term b (or w_0) allows for non-zero threshold



Linearly separable data


- Perceptron works with **linearly separable** data
 - i.e., boundary can be specified by hyperplane
 - E.g., $w_0 + w_1x_1 + \dots + w_kx_k = 0$
- Example: what formula defines the separating hyperplane for these data?



Perceptron: Examples

- A perceptron for AND:
 - Two weights and intercept:
 - $h(x_i) = w_0 + w_1x_{i1} + w_2x_{i2}$
 - One solution:
 - $w_1=0.5, w_2=0.5, w_0=-0.8$

x_1	x_2	y
1	1	T
1	0	F
0	1	F
0	0	F




Note: in drawing these diagrams, we sometimes indicate a threshold T instead of the bias w_0 , such that $T = -w_0$

Perceptron: Your turn

- A perceptron for OR:
 - Two weights and intercept:
 - $h(x_i) = w_0 + w_1x_{i1} + w_2x_{i2}$
 - Find possible weights w_0, w_1, w_2




x_1	x_2	y
1	1	T
1	0	T
0	1	T
0	0	F



Perceptron: Examples

- You've seen AND and OR
- A perceptron for XOR?
- Impossible! → Why?
- XOR is not **linearly separable**

x_1	x_2	y
1	1	F
1	0	T
0	1	T
0	0	F



Outline

- Neural Networks: Motivation and Biology
- The Perceptron
- **Learning perceptron weights**
- Multilayer networks
- Learning multilayer weights
- Backpropagation
- Summary

Learning weights

- Given we have input and output (for instance, a truth table), how do we learn the weights?
- In practice, there are several ways
 - We'll start with Rosenblatt's algorithm (circa 1950's)

Learning weights (Rosenblatt)

- Rosenblatt's Algorithm (perceptron):

```
initialize weights randomly
```

```
while termination condition is not met:
```

```
    initialize  $\Delta w_j = 0$ 
```

```
    for each training example  $(X_i, Y_i)$ :
```

```
        compute predicted output  $\hat{Y}_i$ 
```


```
        foreach weight  $w_j$ :
```

```
             $\Delta w_j = \Delta w_j + \eta (Y_i - \hat{Y}_i) X_i$  ← "error-driven" learning
```

```
    for each weight  $w_j$ :
```

```
         $w_j = w_j + \Delta w_j$  ← Learning rate
```

Perceptron: In action




Who cares?

- Rosenblatt proved the algorithm is guaranteed to converge as long as:
 - Training data are linearly separable
 - Learning rate is sufficiently small
 - (In the proof, it has to be infinitesimally small)

Learning weights: Another Approach

- Output is linear function of weights:
 - (Forget about step function for a moment)
 - $\hat{Y}_i = w_0 + w_1 x_{i1} + \dots + w_n x_{in}$
- Assume error is quadratic function of output
 - $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (Y_i - \hat{Y}_i)^2$
- What does this remind you of?
 - Gradient Descent!
 - $\beta \leftarrow \beta - \frac{R}{N} (Y_i - \hat{Y}_i) X_i$



Training Rule vs. Gradient Descent

- Are these approaches different?
 - Training Rule (Rosenblatt)
 - $\Delta w_j = \Delta w_j + \eta (Y_i - \hat{Y}_i) X_i$
 - Gradient Descent w/ Logistic Regression
 - $\beta \leftarrow \beta + R(Y_i - \hat{Y}_i)X_i$
- The key is the \hat{Y}_i
 - Perceptron: \hat{Y}_i is a step function, either 0 or 1
 - G.D. requires convex surface, not a step function
 - Logit: \hat{Y}_i is a smooth, continuous function

Training Rule vs. Gradient Descent

- Perceptron Training Rule
 - Guaranteed to work if data are linearly separable
 - Requires sufficiently small learning rate η
- Training with Gradient Descent
 - With convex loss...
 - Guaranteed to converge to minimum error
 - Works when data contains noise
 - Works when data are not linearly separable

Perceptron: Summary


- **Online** algorithm: only considers one instance at a time
- **Error-driven**: Only updates on failure
- Guaranteed to converge if solution exists
- But boundary is **linear**

Outline

- Neural Networks: Motivation and Biology
- The Perceptron
- Learning perceptron weights
- **Multilayer networks**
- Learning multilayer weights
- Backpropagation
- Summary

Limitations of the Perceptron



- Only works with linearly separable data
- Only works if learning rate is small enough (Rosenblatt's proof)
- These sort of problems led to "long winter" (1980's)



x_1	x_2	y
1	1	-1
1	0	1
0	1	1
0	0	-1

Multilayer Networks

- Single-layer networks are limited – they can only learn hyperplanes
- What if we layer neurons?
 - Two-layer network
 - (two layers of weights)




Nonlinearity

- OR perceptron:


- $w_1=1, w_2=1, b=-0.9$

x1	x2	z
1	1	1
1	0	1
0	1	1
0	0	-1



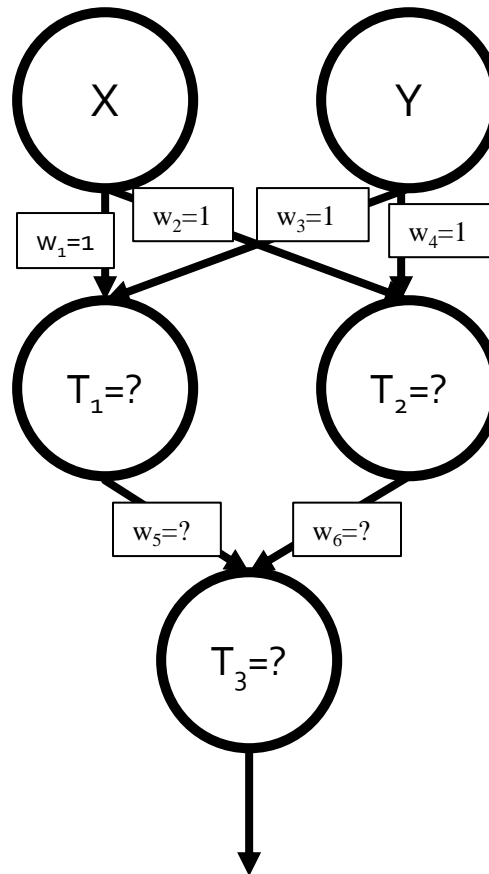
- Two-layer XOR:

x1	x2	z
1	1	-1
1	0	1
0	1	1
0	0	-1



Your Turn: XOR

- What weights complete the XOR MLP?



Universal Approximation Theorem

- Two-Layer Networks are Universal Function Approximators)
 - Let F be a continuous function on a bounded subset of D -dimensional space. Then there exists a two-layer neural network F' with a finite number of hidden units that approximate F arbitrarily well. Namely, for all x in the domain of F ,
$$|F(x) - F'(x)| < \varepsilon$$
- i.e., “two-layer networks can approximate any function”
- But we still might want more than two layers
 - Fewer neurons, time to learn, time to compute, etc.

Universal Approximation Theorem

- This is a powerful theorem, but...
 - “Just because a function can be represented does not mean it can be learned”
- Learning may require:
 - Insane complexity
 - Insane amounts of data
 - Insane computational resources

Outline

- Neural Networks: Motivation and Biology
- The Perceptron
- Learning perceptron weights
- Multilayer networks
- **Learning multilayer weights**
- Backpropagation
- Summary