

THIS IS YOUR MACHINE LEARNING SYSTEM?

YUP! YOU POUR THE DATA INTO THIS BIG PILE OF LINEAR ALGEBRA, THEN COLLECT THE ANSWERS ON THE OTHER SIDE.

WHAT IF THE ANSWERS ARE WRONG?

JUST STIR THE PILE UNTIL THEY START LOOKING RIGHT.



INFO 251: Applied Machine Learning

Neural Networks, part 2

Announcements

- PS₄ due next Monday
- PS₅ (Trees, forests, basic neural nets) posted next week

Key concepts (last class)

- Mimicking basic neural processes
- The perceptron
- Perceptron limitations
- Rosenblatt's algorithm
- Perceptron training vs. gradient descent
- Multilayer networks
- Universal approximation theorem

Outline

- **Learning multilayer weights: Intuition**
- Generalizing logistic regression
- Learning multilayer weights: simple case
- Backpropagation: Intuition
- Backpropagation: Video
- Summary

Multi-layer networks are great!

- Multi-layer networks have great properties, e.g. can solve the XOR problem – this was recognized by Minsky and Papert (1969)
- However, people didn't know how to fit these multi-layer networks. What weights?
 - The perceptron training rule we discussed doesn't work with multiple layers
 - Still, no one has figured out how to generalize Rosenblatt's algorithm to multi-layer networks



Intuition check 1

- Can off-the-shelf gradient descent, such as what you're implementing in PS4, be used to learn the weights in multi-layer networks?
 - No! The “activation” step function of a perceptron creates non-convex loss. Not differentiable

$$\hat{Y}_i = \begin{cases} 1 & \text{if } w_0 + w_1x_{i1} + \dots + w_nx_{in} > 0 \\ -1 & \text{otherwise} \end{cases}$$

Intuition check 2

- Well, if the issue is with the step function, can we just omit the step function entirely?
- No!
 - Without activation, each individual unit is linear
 - $\hat{Y}_i = w_0 + w_1x_{i1} + \dots + w_nx_{in}$
- Combination of units is more complex...
 - $\hat{Y}_i = v_0 + v_1(w_0 + w_1x_{i1} + \dots + w_mx_{im}) + \dots + v_nx_{in}$
 - But still linear!




So, what do we do?

- We need something like a step function to capture nonlinearities

$$\hat{y}_i = \begin{cases} 1 & \text{if } w_0 + w_1 x_{i1} + \dots + w_n x_{in} > 0 \\ -1 & \text{otherwise} \end{cases}$$

- But the step function itself creates issues for learning weights
 - It's nonlinear (good!), but not differentiable (bad!)
 - Gradient descent needs a differentiable function
- In other words, we need a nonlinear, differentiable function
 - Examples: sigmoid function, tanh()




Outline

- Learning multilayer weights: Intuition
- **Generalizing logistic regression**
- Learning multilayer weights: simple case
- Backpropagation: Intuition
- Backpropagation: Video
- Summary

Generalizing logistic regression

- Recall the loss functions from before
 - Linear Regression
 - $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (Y_i - \alpha - \beta X_i)^2 [+ \lambda \sum_{j=1}^k \theta_j^2]$
 - Logistic Regression (omitting regularization)
 - $J(\theta) = -\frac{1}{N} \sum_{i=1}^N Y_i \cdot \log \hat{Y}_i + (1 - Y_i) \log(1 - \hat{Y}_i)$
 - Note: Something like this could work on the **last** (output) layer of a multi-layer network, to determine the weights on that layer
 - Such as the bold \mathbf{v}_i 's in the diagram



Generalizing logistic regression

- However, neural networks have multiple layers/outputs (generalization of logistic function)

$$J(\theta) = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K Y_{ik} \cdot \log \hat{Y}_{ik} + (1 - Y_{ik}) \log(1 - \hat{Y}_{ik})$$

- (This is just the cost function for a network with k outputs)
 - Here, each \hat{Y}_{ik} is like a nested set of logistic regressions
- BUT: we don't have the target value (the Y_{ik}) for hidden layers!

Outline


- Learning multilayer weights: Intuition
- Generalizing logistic regression
- **Learning multilayer weights: simple case**
- Backpropagation: Intuition
- Backpropagation: Video
- Summary

Two-layer backpropagation

- Backprop = gradient descent + chain rule
 - Solution and algorithm for two layers not too complex (see Daume)

■ Objective:
$$\min_{\mathbf{W}, \mathbf{v}} \sum_n \frac{1}{2} \left(y_n - \underbrace{\sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}_n)}_{\text{This is the prediction } \hat{y}_n} \right)^2$$

- n indexes observations
- i indexes hidden units
- \mathbf{v} is second layer weights
- f is the sigmoid function
- \mathbf{w}_i is the vector of weights feeding into node i
- \mathbf{W} is first layer weights



Two-layer backpropagation

- Objective:

$$\min_{\mathbf{W}, \mathbf{v}} \sum_n \frac{1}{2} \left(y_n - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}_n) \right)^2$$

- Apply chain rule:

$$\mathcal{L}(\mathbf{W}) = \frac{1}{2} \left(y - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}) \right)^2$$


$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}_i} = \frac{\partial \mathcal{L}}{\partial f_i} \frac{\partial f_i}{\partial \mathbf{w}_i}$$

$$\frac{\partial \mathcal{L}}{\partial f_i} = - \left(y - \sum_i v_i f(\mathbf{w}_i \cdot \mathbf{x}) \right) v_i = -e v_i$$

$$\frac{\partial f_i}{\partial \mathbf{w}_i} = f'(\mathbf{w}_i \cdot \mathbf{x}) \mathbf{x}$$

- Solution:

$$\nabla_{\mathbf{w}_i} = -e v_i f'(\mathbf{w}_i \cdot \mathbf{x}) \mathbf{x}$$



prediction error ($y - \hat{y}$)

Learning multilayer weights

- Solution with two layers

$$\nabla_{w_i} = -ev_i f'(w_i \cdot x)x$$

- Does this make sense?
 - If predictive error (e) is small, take small steps
 - If v_i is small, hidden unit i has little influence on output, gradient should be small
 - If e or v_i changes sign, gradient should also flip sign

Outline

- Learning multilayer weights: Intuition
- Generalizing logistic regression
- Learning multilayer weights: simple case
- **Backpropagation: Intuition**
- Backpropagation: Video
- Summary


Backpropagation

- How to generalize from two layers?
- Sketch of procedure
 1. Forward Propagation -> Outputs
 2. Backward Propagation -> Generate “deltas”
 3. Weight Update -> same as in gradient descent

Intuition

■ Forward Propagation


- Given a training example (X_1, \dots, X_n) and output Y_i :
- Propagate inputs/activations forward, applying sigmoid function on dot products



Intuition

■ Backward Propagation


- For a single training example i :
 - $\text{Cost}(i) = Y \cdot \log \hat{Y}_i + (1 - Y_i) \log(1 - \hat{Y}_i)$
 - i.e., how close is output to actual value?
- Idea is to propagate costs backwards to earlier nodes
 - Compute $\delta_{jK} = \text{"error" of } j^{\text{th}} \text{ node in } K^{\text{th}} \text{ (output) layer}$
 - $\delta_{jK} = Y_{jK}(1 - Y_{jK})(\hat{Y}_{jK} - Y_{jK})$
 - Note: getting these partials is a bit complex, but mostly just chain rule + gradient descent (see Hastie ch. 11)



Intuition

■ Update weights


- For each hidden unit h in k^{th} layer, update each weight as $+\eta\delta_{hk}x_i$




Outline

- Learning multilayer weights
- Generalizing logistic regression
- Learning multilayer weights: simple case
- Backpropagation: Intuition
- **Backpropagation: Video**
- Summary

What is back-propagation?



Video #2



Outline

- Learning multilayer weights
- Generalizing logistic regression
- Learning multilayer weights: simple case
- Backpropagation: Intuition
- Backpropagation: Video
- **Summary**

Neural Networks: Issues

- Non-convex, sensitive to initialization
 - Common solutions: Randomize initialization (small random/uniform weights), train multiple networks
- Avoiding overfitting
 - Early stopping
 - Penalize large weights (explicit regularization)
 - Include fewer layers, weights per layer

Tuning networks

- Many considerations
 - How many layers?
 - How many units per layer?
 - How to initialize?
 - What learning rate?
 - Weight regularization?
 - When to stop?
- Tuning deep networks can take time
 - Network architecture
 - Layer-wise initialization
 - Alternative optimization

Neural Networks: Summary

- Very flexible, can model complex and non-linear relationships
- Compute-intensive
- Can be parallelized!
- Very hard to interpret, i.e. “black box”